TET NETWORK

A Fluid Peer-to-Peer Computational and Energy Resource Protocol

The AI-Native Sovereign Layer 1

Version	**Genesis Draft v1.0**
Date	**April 28, 2026**
Author	**Steve**
Role	**Founder-Architect, TET Network Project**
Contact	**yizhenxianshi@gmail.com**

1. FOREWORD

Legacy cloud infrastructure is not an immutable reality. It is a historical accident—a temporary monopoly over human intelligence by capital, sustained by switching costs and institutional inertia, not by any underlying technical superiority.

The contradiction is structural: AI, the most consequential technology of this era, runs almost entirely on compute owned by four companies. Those companies set the price of intelligence. That is not a market. It is a tollbooth.

TET Network is the architectural response. Its central premise—Compute is Energy—is not new. What is new is the conclusion drawn from it: if compute and energy are fungible, then any device that consumes electricity is a potential mining node, a potential inference provider, a participant in the sovereign compute grid. Traditional chains impose a single consensus rule on every participant. A smartphone competing against an A100 cluster under identical rules is not decentralization; it is a performance of decentralization that produces the same capital-intensive oligopoly as before.

Context-Aware Adaptive Consensus (CAAC) eliminates that constraint. The chain adapts to the hardware. A GPU farm runs Proof of Compute. A smartphone runs Proof of Relay. The protocol unifies them under a single economic model without demanding identical capability.

This paper defines the complete architecture. Implementation requires engineers who can operate at the intersection of distributed systems, applied cryptography, and ML infrastructure. If the design contains a flaw, that flaw should be found in a technical critique, not in a deployed mainnet.

2. ABSTRACT & EXECUTIVE SUMMARY

AI infrastructure currently runs on a trust model that concentrates risk: a small number of platform operators control access, pricing, and availability. The consequences—opaque censorship, high intermediary margins, correlated outage risk—are not incidental. They are structurally inevitable in any centralized system.

TET Network is a Post-Quantum, AI-Native Sovereign Layer 1 designed to replace that trust model with mathematical guarantees. Four properties define the architecture:

— **Context-Aware Adaptive Consensus (CAAC):** The protocol dynamically assigns each node a validation role proportional to its hardware profile. High-end GPU clusters execute Proof of Compute; edge devices execute Proof of Relay. Hardware heterogeneity is a feature, not a flaw.

— **Universal Miner Inclusion:** By adapting consensus to the environment rather than the reverse, TET eliminates the hardware barrier that confines PoW and PoS participation to well-capitalized actors. Any device with a network connection and a power source is a valid participant.

— **Thermodynamic Efficiency:** No computation cycle in TET is expended on abstract puzzles. Every watt processed by a PoC node resolves a real AI inference task. Energy expenditure is directly proportional to economic output.

— **Post-Quantum Security:** ML-DSA (FIPS 204 / Dilithium) is implemented as the base-layer signature scheme from genesis, providing mathematical resistance to Shor's algorithm without a migration event.

3. THE STRUCTURAL FAILURE OF HOMOGENEOUS CONSENSUS

Every major Layer 1 protocol in production today assumes that all nodes compete under identical rules. In Proof of Work, this produces ASIC centralization: the hardware arms race converges toward a small number of industrial operators who can amortize fabrication costs. Proof of Stake substitutes capital for hardware but preserves the monopoly structure—the validator set is bounded by who can afford to lock the minimum stake.

Neither model wastes compute on external utility. Bitcoin's SHA-256 hashing, Ethereum's former Ethash, and their successors produce cryptographic proofs that secure the ledger and nothing else. The energy expenditure is real; the economic output beyond consensus is zero.

This paper terms the pattern Consensus Decay: the progressive concentration of network control into a small actor class, regardless of whether the scarce resource is hardware or capital. TET addresses Consensus Decay at the protocol level by decoupling the consensus mechanism from any single hardware requirement and routing the underlying computational work toward externally useful AI inference.

4. CORE INNOVATION: CONTEXT-AWARE ADAPTIVE CONSENSUS (CAAC)

CAAC is a protocol-level mechanism that evaluates the physical capabilities and network context of a connecting node and assigns it an optimal validation role. The design principle is direct: the chain adapts to the hardware, not the reverse.

4.1 High-Performance Nodes — Proof of Compute (PoC)

Nodes with substantial processing capacity—data-center GPU clusters being the canonical case—are assigned the PoC role. These nodes process real AI inference tasks: matrix operations, attention computations, token generation. The network verifies output through optimistic execution backed by Zero-Knowledge fraud proofs. TET rewards are issued in direct proportion to verified thermodynamic energy provided.

4.2 Edge Nodes — Proof of Relay (PoR) and PQC Validation

Constrained devices—smartphones, IoT hardware—are assigned the PoR role. They do not run heavy inference. They sustain the network by routing data packets, maintaining state telemetry, and executing lightweight verification of Post-Quantum Cryptographic signatures. The economic reward is calibrated to the work performed: smaller than PoC, but non-trivial, and accessible to any device with a data connection.

The two roles share a single ledger, a single token, and a single security model. CAAC is the routing layer that connects them.

5. THE FLUID ARCHITECTURE AND SOVEREIGN RUNTIME

TET Network instantiates the CAAC framework as the Environment-Adaptive Chain—the Fluid Chain. The network eliminates abstract benchmarking scores. TET is the universal unit of both value and physical resource: 1 TET represents a deterministic quantity of verified AI computation or network sustainment work, enforced at the protocol level.

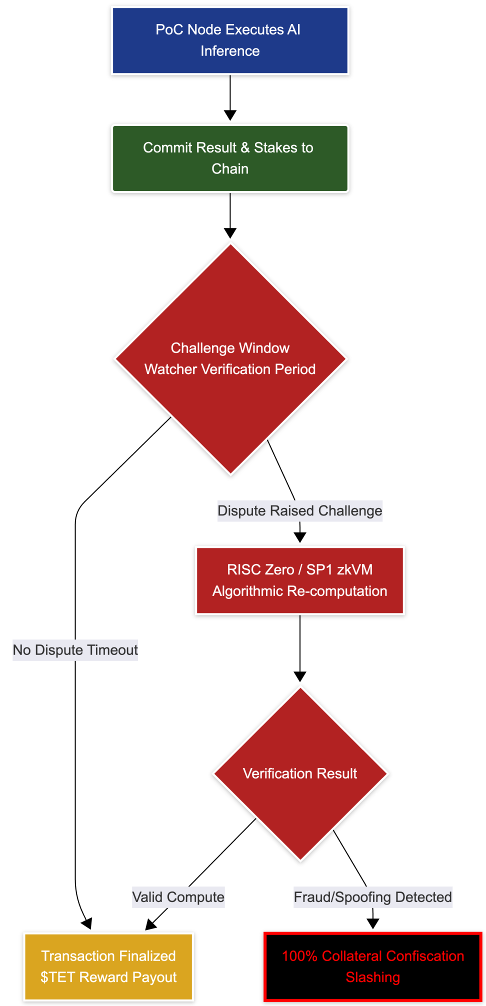

5.1 The Sovereign Runtime

Ethereum meters compute steps using Gas, an abstract unit that prices execution without reference to external utility. TET prices execution in units of thermodynamic work directly. When a smart contract requests an AI inference, the Sovereign Runtime routes the task to a PoC node via an Optimistic Verification model: the node executes off-main-thread and returns a cryptographic commitment to the result. The main chain accepts the commitment optimistically during the challenge window. Disputes invoke the ZK-Court.

The ZK-Court uses RISC Zero or SP1 zkVMs to replay the disputed execution and produce a verifiable proof. A commitment that fails verification triggers a 100% slash of the malicious node's staked TET. The cost of fraud is not a fine; it is confiscation.

5.2 Mathematical Model of Thermodynamic Work

The value of 1 TET is not arbitrary. It is a cryptographic proof of physical work. The expected reward R for a node over interval T is:

$$R(T) = \frac{\sum_{i \in \texttt{verified\_tasks}(T)} \left[\eta(W_i) \cdot C(t_i)\right]}{D(t)}$$

where $\eta(W_i)$ is the verified thermodynamic output of task i, $C(t_i)$ is the network's compute price at time t_i, and D(t) is the dynamic difficulty factor. This equation establishes the sovereign peg: intelligence generation cryptographically bound to physical electricity.

6. ANATOMY OF A FLUID TRANSACTION

A TET transaction represents either a value transfer or a physical AI computation request. The transaction structure carries a Workload Flag—a single binary field—that instructs CAAC on routing:

— **Flag = 0:** Standard value transfer. Validated by edge nodes via PoR. Low-latency, minimal resource requirement.

— **Flag = 1:** AI inference or compute request. Routed exclusively to PoC high-performance nodes. Subject to ZK-Court dispute resolution.

This distinction is enforced at the protocol layer. A Flag-1 transaction submitted to an edge node is rejected by the routing logic; it cannot be processed at reduced cost by misrouting.

7. EDGE CLUSTERING AND SEMANTIC ROUTING

Geographic distribution of compute provides censorship resistance but introduces latency. Distributing a 50GB model weight file to every node that might receive an inference request is not viable. TET resolves this through Weight Locality.

Nodes specialize and cache specific model families according to their hardware profile. When a user submits an inference request, the protocol routes the task to the nearest logical cluster that already holds the required model weights in memory. Data transfer overhead is reduced to the inference input and output, not the model itself. Decentralized inference latency converges toward centralized cloud latency without sacrificing the distributed trust model.

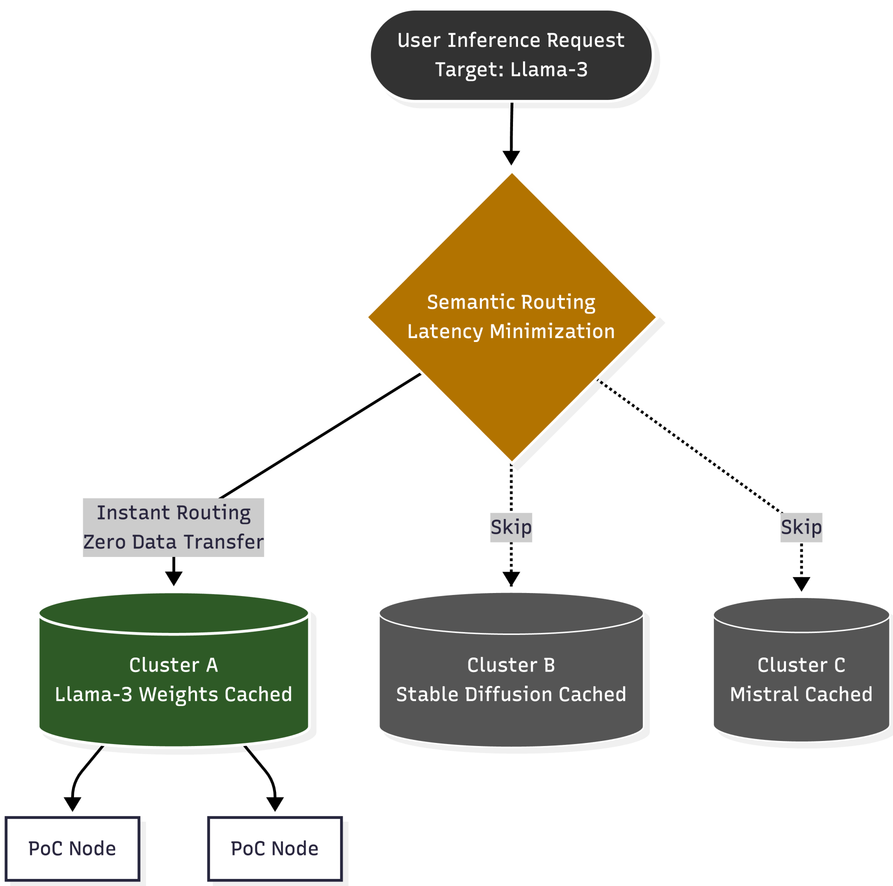

8. STATE VERIFICATION FOR EDGE DEVICES

The 8 Billion Device Plan—the goal of making every internet-connected device a valid network participant—requires that smartphones and IoT hardware can participate securely without storing multi-terabyte ledger history. TET addresses this through a multi-level Merkle tree structure.

PoC nodes store the complete state. PoR edge nodes operate as light clients: they download block headers only and verify the cryptographic proofs attached to them. To query a specific contract state or balance, an edge device retrieves only the relevant Merkle branch. The proof is mathematically sufficient to confirm data validity without the full chain. Storage requirements for an edge node are bounded by the header chain and a small local proof cache—well within the constraints of consumer hardware.

9. THERMODYNAMIC EFFICIENCY AND NETWORK RESILIENCE

TET implements a zero-waste computational model. Every unit of energy consumed by a PoC node is allocated to an externally useful task. There are no abstract hash puzzles. The protocol does not burn electricity to prove that electricity was burned.

Network resilience is a structural property of the two-tier architecture, not an engineered failsafe. If major data centers or GPU clusters go offline—due to hardware failure, regulatory action, or geopolitical event—the surrounding PoR edge nodes automatically sustain state propagation and routing. The network degrades in inference throughput but does not lose ledger continuity. Millions of independently operated, geographically distributed light nodes provide redundancy that no centralized data center can match.

This property is sometimes described internally as the Cockroach Doctrine: the network cannot be physically killed because it has no single physical location to kill.

10. POST-QUANTUM SECURITY

ECDSA—the signature scheme protecting the majority of current blockchain assets—is vulnerable to Shor's algorithm on a sufficiently capable quantum computer. The timeline for that capability is uncertain; the mathematical vulnerability is not. Migrating a deployed mainnet to a new signature scheme after the fact is a coordination problem with no clean solution.

TET implements ML-DSA (Module-Lattice-Based Digital Signature Algorithm, FIPS 204, standardized by NIST in 2024) as the base-layer signature scheme from genesis block zero. Private keys and on-chain assets are mathematically resistant to both classical and quantum adversaries from day one. There is no migration event, no opt-in period, no legacy ECDSA compatibility layer that introduces attack surface.

11. TOKENOMICS: GENESIS ALLOCATION AND SCARCITY

Maximum total supply is fixed at 10,000,000,000 TET. The allocation is hard-coded into the base layer and is not subject to governance modification.

11.1 Genesis Allocation

— 25% — Founder & Core Contributors: Subject to a multi-year vesting schedule. Not liquid at genesis.

— 50% — Resource Mining Emission: Programmatically emitted over decades to reward PoC and PoR nodes. Released algorithmically; not controlled by any entity.

— 25% — Ecosystem Treasury: Held in decentralized smart contracts. Funds AI infrastructure grants and network expansion via on-chain governance.

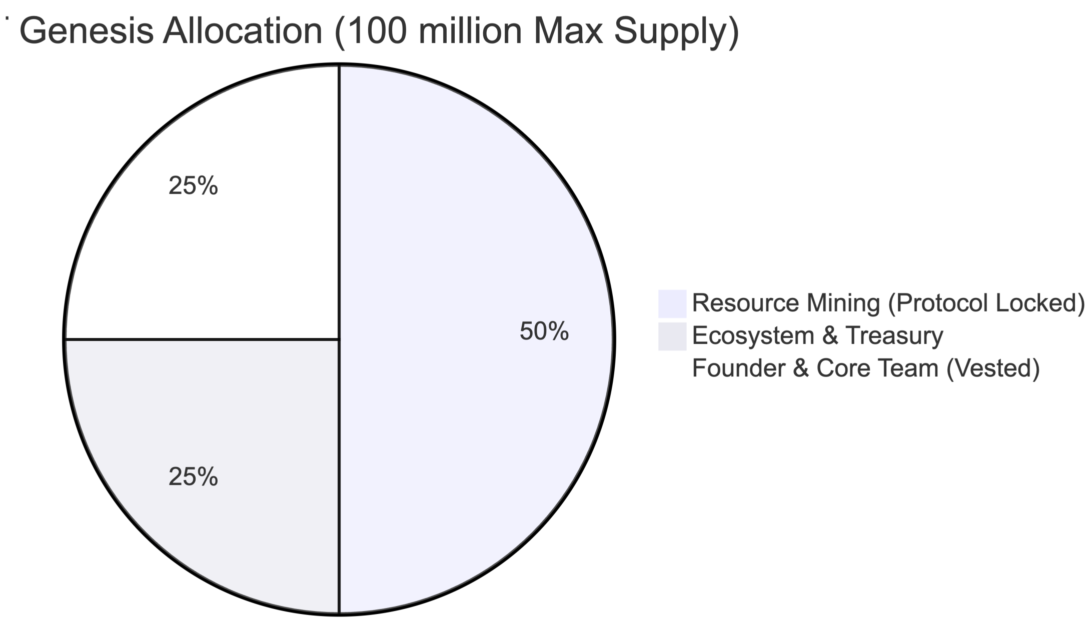

11.2 Deflationary Burn Mechanism

50% of all transaction fees—including AI inference costs—are permanently burned on settlement. As network demand grows, the burn rate increases. The supply curve is structurally deflationary under sustained usage, without requiring any protocol intervention.

12. APPLICATIONS OF THE SOVEREIGN FLUID GRID

Ethereum generalized state transitions and made DeFi possible. TET generalizes thermodynamic compute. The application surface extends beyond decentralized inference to any system where autonomous, verifiable AI execution is a primitive.

12.1 Rollup-as-a-Service with Native AI (RaaS)

Deploying a dedicated blockchain today requires significant engineering resources and ongoing operational overhead. TET removes that barrier. A single transaction deploys a sovereign Layer-2 sub-chain anchored to TET-Core. These sub-chains are not ledgers with a computation layer bolted on—they are AI-native by construction. The base layer's PoC network provides each sub-chain with a "cannot be evil" AI monitoring layer: autonomous smart-contract auditing, anomaly detection, and block-space optimization running at the infrastructure level, not the application level.

Deployment and operational costs are denominated exclusively in TET, creating a persistent demand sink for the native asset regardless of application-layer activity.

12.2 Decentralized AI Inference Markets

Startups and researchers can route LLM queries through TET at a fraction of centralized cloud costs. Every inference result carries a cryptographic proof of correct execution, verifiable via the ZK-Court without trusting the executing node.

12.3 Autonomous AI Agents — DAOs 2.0

Smart contracts can natively invoke PoC nodes to run localized AI models—analyzing on-chain data, executing trades, initiating governance proposals—without human intervention at each decision point. The AI execution layer is not an external oracle; it is part of the consensus model.

12.4 Quantum-Secure Financial Primitives

Any DeFi application deployed on TET inherits ML-DSA security without additional configuration. As quantum capability advances, TET-native financial contracts remain mathematically sound while ECDSA-based chains face an existential migration event.

12.5 Neural State Transition — The World Brain

During idle periods, PoC nodes contribute spare compute to continuous federated fine-tuning of a globally shared, open-source base model. The cryptographic state of the chain becomes a living neural network: tamper-evident, censorship-resistant, and continuously refined by privacy-preserving telemetry from edge nodes worldwide. The model belongs to the protocol. No corporation can revoke access to it.

12.6 Sentient Assets — Smart Contracts 2.0

With the Neural State Transition integrated into the runtime, contract logic is no longer restricted to deterministic if-then rules. Developers can deploy assets with embedded inference: sovereign wallets that analyze transaction context to block social-engineering attacks; digital entities that negotiate their own market value based on real-time ecosystem data. The asset itself reasons about its environment.

13. DEVELOPMENT ROADMAP

— Phase 0 — The Inference Wedge: Launch of a lean, highly optimized LLM inference market targeting edge nodes. The protocol establishes the baseline energy-to-compute ratio using a standardized Llama-3 model. No consensus layer yet; this phase validates the economic model.

— Phase 1 — CAAC Implementation: Full deployment of Context-Aware Adaptive Consensus. The network begins autonomously categorizing nodes into PoC and PoR roles based on real-time hardware telemetry. ZK-Court dispute resolution is activated.

— Phase 2 — Post-Quantum Fluid Grid: Complete integration of ML-DSA signatures. The network achieves full fluid distribution: AI requests routed dynamically across a global, quantum-resistant matrix of up to 8 billion devices. Federated learning and Sentient Asset primitives go live.

14. THREAT MODELS AND CRYPTOGRAPHIC DEFENSES

14.1 Lazy Evaluation — Compute Spoofing

A malicious PoC node may attempt to conserve electricity by submitting a fabricated inference result rather than executing the task. The ZK-Court defends against this. During the challenge window, any Watcher can trigger a RISC Zero or SP1 zkVM execution trace against the node's submitted commitment. A trace that contradicts the commitment triggers a 100% slash of the node's staked TET. The attack has a fixed, non-recoverable cost; the expected value is permanently negative.

14.2 Hardware Spoofing — Sybil Attack

A node may attempt to emulate a high-end GPU using lower-tier hardware to claim inflated PoC rewards. TET mitigates this via Probabilistic Hardware Fingerprinting: the protocol issues non-deterministic micro-tasks that exploit specific execution timing characteristics of real hardware. The response profile is an unforgeable cryptographic proof of the physical hardware layer. Emulation cannot replicate the timing signatures of genuine silicon.

14.3 The Mathematics of Economic Finality

TET enforces a Cost-of-Corruption model in which the economic cost of any Byzantine attack strictly exceeds its potential profit. The slashed collateral S is defined as:

$$S = \lambda \cdot R_{\texttt{expected}}$$

Because S > R_expected for any feasible attack, the expected value of fraud is permanently locked negative. The protocol does not require trusting nodes; it requires only that nodes behave as rational economic actors—a weaker and more reliable assumption.

15. CONCLUSION

The structural flaws of contemporary Layer 1 architectures are not implementation bugs. They are design choices: homogeneous consensus that concentrates power, energy expenditure with no external utility, cryptographic assumptions that will not survive quantum hardware, and an implicit requirement that meaningful participation demands significant capital.

TET Network addresses each of these at the protocol level. Context-Aware Adaptive Consensus dissolves the hardware monopoly. The Sovereign Runtime converts energy expenditure into economically useful AI inference. ML-DSA provides quantum-resistant security from genesis. The two-tier node model extends participation to any device on Earth.

The result is not an incremental improvement over existing chains. It is a different kind of infrastructure: fluid, autonomous, and sovereign—where artificial intelligence compute and physical energy are treated as universal resources, mathematically secured and accessible without permission.

16. CALL TO BUILDERS

The architecture is complete. Translating it into a production codebase is an engineering problem, not a research problem—but it is not a trivial one. We are looking for Core Contributors with demonstrated expertise in:

— **Systems Programming:** Rust. Low-level networking, memory-safety-critical consensus code.

— **Decentralized Networking:** libp2p. Peer discovery, NAT traversal, protocol multiplexing at scale.

— **Applied Cryptography:** zkVM development (RISC Zero, SP1), ML-DSA implementation, ZK fraud proof construction.

— **Distributed ML:** Federated learning, quantized inference, model weight distribution across heterogeneous hardware.

We do not want résumés. If this architecture contains an inefficiency or a superior mechanism exists, locate it, describe it precisely, and send a technical critique to:

```
yizhenxianshi@gmail.com    [Subject: Core Builder Application]
```

The first block of the global resource grid will not be mined by credential. It will be mined by those who understand why the architecture is correct and can prove where it is not.

Tet-core v0.1 ↓

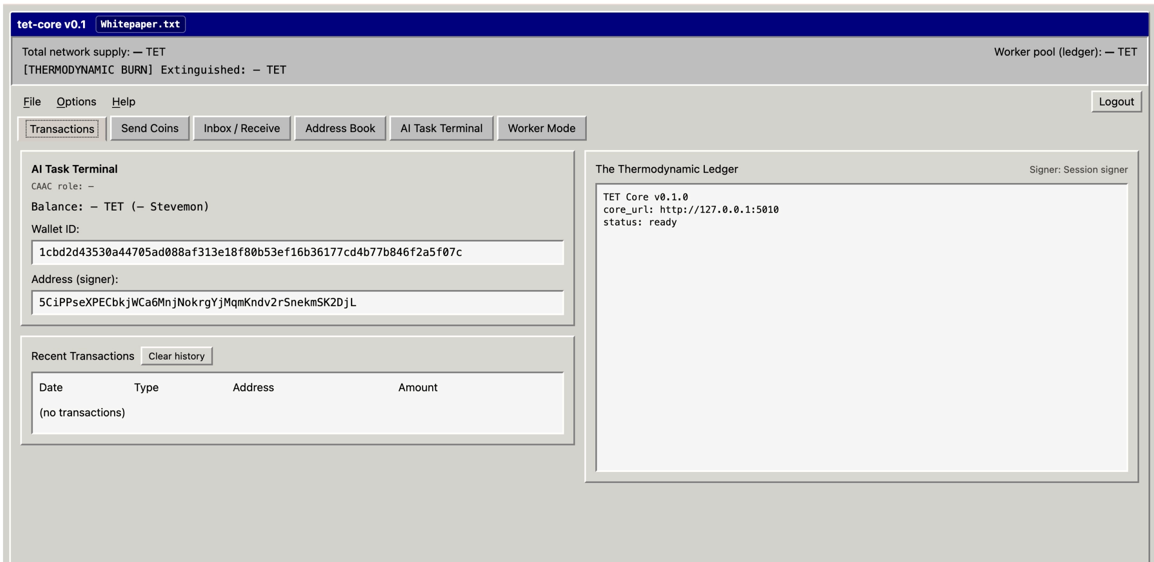
tet-core v0.1
Whitepaper.txt
Total network supply: — TET
[THERMODYNAMIC BURN] Extinguished: — TET
Worker pool (ledger): — TET
File
Options
Help
Logout
Transactions
Send Coins
Inbox / Receive
Address Book
AI Task Terminal
Worker Mode
AI Task Terminal
CAAC role: —
Balance: — TET (— Stevemon)
Wallet ID:
1cbd2d43530a44705ad088af313e18f80b53ef16b36177cd4b77b846f2a5f07c
Address (signer):
5CiPPseXPECbkjWCa6MnjNokrgYjMqmKndv2rSnekmSK2DjL
Recent Transactions
Clear history
Date
Type
Address
Amount
(no transactions)
The Thermodynamic Ledger
Signer: Session signer
TET Core v0.1.0
core_url: http://127.0.0.1:5010
status: ready

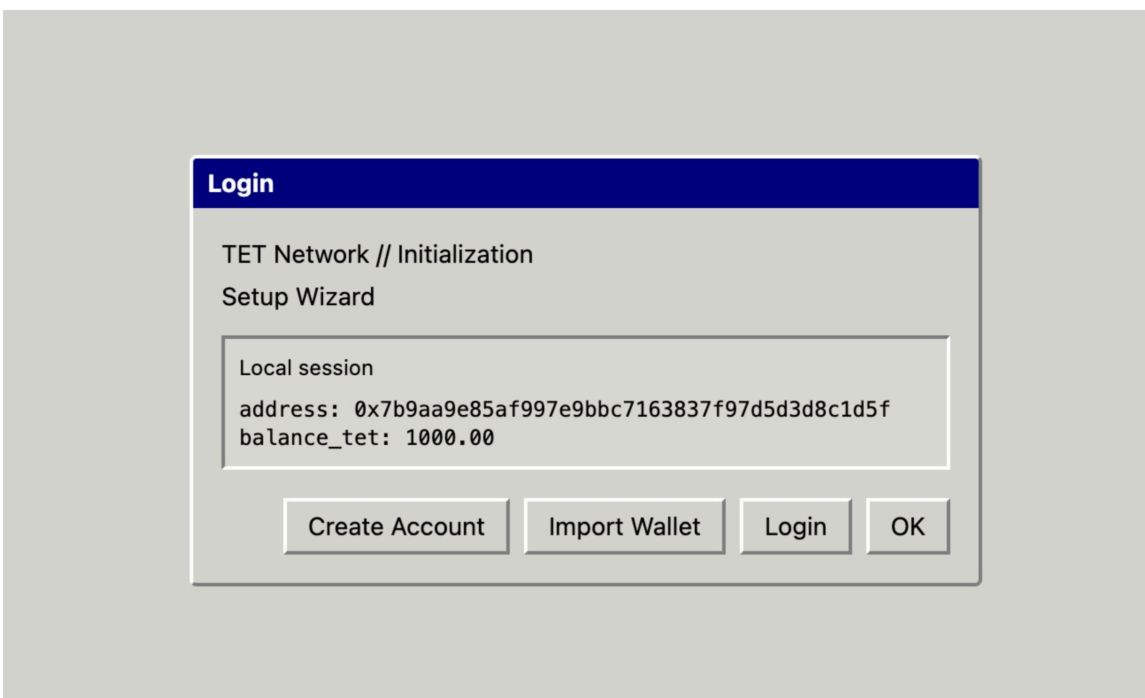
Login
TET Network // Initialization
Setup Wizard
Local session
address: 0x7b9aa9e85af997e9bbc7163837f97d5d3d8c1d5f
balance_tet: 1000.00
Create Account
Import Wallet
Login
OK

17. REFERENCES

Core Consensus & Cryptoeconomics

[1] Nakamoto, S. (2008). Bitcoin: A Peer-to-Peer Electronic Cash System.

[2] Buterin, V. (2014). Ethereum: A Next-Generation Smart Contract and Decentralized Application Platform.

[3] Sompolinsky, Y., & Zohar, A. (2015). Secure High-Rate Transaction Processing in Bitcoin. Financial Cryptography and Data Security.

[4] Buterin, V., & Griffith, V. (2017). Casper the Friendly Finality Gadget. Ethereum Foundation.

Zero-Knowledge & Post-Quantum Cryptography

[5] Ben-Sasson, E., Chiesa, A., et al. (2014). SNARKs for C: Verifying Program Executions Succinctly and in Zero Knowledge. CRYPTO 2013.

[6] RISC Zero Team. (2023). RISC Zero: A Zero-Knowledge Virtual Machine for General Purpose Computing.

[7] Succinct Labs. (2024). SP1: A Next-Generation zkVM for High-Performance Proof Generation.

[8] NIST. (2024). FIPS 204: Module-Lattice-Based Digital Signature Standard (ML-DSA).

Decentralized AI & Network Layer

[9] Borzunov, A., et al. (2022). Petals: Collaborative Inference and Fine-tuning of Large Models. arXiv:2209.01188.

[10] Rao, Y. (2021). Bittensor: A Peer-to-Peer Intelligence Market.

[11] AI@Meta. (2024). Llama 3 Model Card. Meta AI Research.

[12] Protocol Labs. (2019). libp2p: A Modular Network Stack for Peer-to-Peer Applications.

Thermodynamics & Hardware Security

[13] McMahan, B., et al. (2017). Communication-Efficient Learning of Deep Networks from Decentralized Data. AISTATS 2017.

[14] Landauer, R. (1961). Irreversibility and Heat Generation in the Computing Process. IBM Journal of Research and Development.

[15] Bennett, C. H. (1982). The Thermodynamics of Computation.

[16] Suh, G. E., & Devadas, S. (2007). Physical Unclonable Functions for Device Authentication and Secret Key Generation. DAC 2007.